


Ensemble learning, bias-variance decomposition

Victor Kitov

v.v.kitov@yandex.ru



Loss vs. model complexity



Comments:

- expected loss on test set is always higher than on train set.
- left to A: model too simple, underfitting, high bias
- right to A: model too complex, overfitting, high variance

Table of Contents

- 1 Bias-variance decomposition
- 2 Ensemble learning
- 3 Fixed integration schemes for classification
- 4 Stacking
- 5 Train set sampling ensemble methods

Bias-variance decomposition

- True relationship $y = f(x) + \varepsilon$
- This relationship is estimated using random training set $(X, Y) = \{(x_n, y_n), n = 1, 2 \dots N\}$
- Recovered relationship $\hat{f}(x)$, x -some fixed constant
- Noise ε is independent of any X, Y , $\mathbb{E}\varepsilon = 0$

Bias-variance decomposition

$$\begin{aligned} \mathbb{E}_{X,Y,\varepsilon}\{[\hat{f}(x) - y(x)]^2\} &= \left(\mathbb{E}_{X,Y}\{\hat{f}(x)\} - f(x)\right)^2 \\ &\quad + \mathbb{E}_{X,Y}\left\{[\hat{f}(x) - \mathbb{E}_{X,Y}\hat{f}(x)]^2\right\} + \mathbb{E}\varepsilon^2 \end{aligned}$$

- Intuition: $MSE = \text{bias}^2 + \text{variance} + \text{irreducible error}$
 - darts intuition

Proof of bias-variance decomposition

Define for brevity of notation $f = f(x)$, $\hat{f} = \hat{f}(x)$, $\mathbb{E} = \mathbb{E}_{X,Y,\varepsilon}$.

$$\begin{aligned}\mathbb{E}(\hat{f} - f)^2 &= \mathbb{E}(\hat{f} - \mathbb{E}\hat{f} + \mathbb{E}\hat{f} - f)^2 = \mathbb{E}(\hat{f} - \mathbb{E}\hat{f})^2 + (\mathbb{E}\hat{f} - f)^2 \\ &\quad + 2\mathbb{E}[(\hat{f} - \mathbb{E}\hat{f})(\mathbb{E}\hat{f} - f)] \\ &= \mathbb{E}(\hat{f} - \mathbb{E}\hat{f})^2 + (\mathbb{E}\hat{f} - f)^2\end{aligned}$$

We used that $(\mathbb{E}\hat{f} - f)$ is a constant w.r.t. X, Y and hence

$$\mathbb{E}[(\hat{f} - \mathbb{E}\hat{f})(\mathbb{E}\hat{f} - f)] = (\mathbb{E}\hat{f} - f)\mathbb{E}(\hat{f} - \mathbb{E}\hat{f}) = 0.$$

$$\begin{aligned}\mathbb{E}(\hat{f} - y)^2 &= \mathbb{E}(\hat{f} - f - \varepsilon)^2 = \mathbb{E}(\hat{f} - f)^2 + \mathbb{E}\varepsilon^2 - 2\mathbb{E}[(\hat{f} - f)\varepsilon] \\ &= \mathbb{E}(\hat{f} - \mathbb{E}\hat{f})^2 + (\mathbb{E}\hat{f} - f)^2 + \Delta\end{aligned}$$

Here $\mathbb{E}[(\hat{f} - f)\varepsilon] = \mathbb{E}[(\hat{f} - f)]\mathbb{E}\varepsilon = 0$ since ε is independent of X, Y .

Table of Contents

- 1 Bias-variance decomposition
- 2 Ensemble learning
 - Ensemble learning use cases
- 3 Fixed integration schemes for classification
- 4 Stacking
- 5 Train set sampling ensemble methods


Ensemble learning

- Ensemble model - model using predictions of other models.
- Example: stacking
 - suppose we have base models $\hat{y}_1 = f_1(x), \dots, \hat{y}_M = f_M(x)$.
 - stacking: $\hat{y}(x) = G(f_1(x), \dots, f_M(x))$
- Used in
 - supervised methods: regression, classification, collaborative filtering.
 - unsupervised methods: clustering, dimensionality reduction.

- 2 Ensemble learning
 - Ensemble learning use cases

Multiclass classification using binary classifiers


Multiclass classification with one-vs-rest, one-vs-one, error correcting codes schemes:



Solve underfitting


- Suppose $f_1(x), \dots, f_M(x)$ are too simple and underfit.
- May increase complexity by applying $G(f_1(x), \dots, f_M(x))$

Example



Example

Classifier 1 and classifier 2 combined using AND rule




Solve overfitting

- $f_1(x), \dots, f_M(x)$ overfit (have high variance)
 - decision trees on different training sets
 - neural networks estimated with different initial conditions
- Regression: average their variability to get more robust estimate:

$$\hat{y}(x) = \frac{1}{M} \sum_{m=1}^M f_m(x)$$


- Classification: majority voting.

Regression: high variance




Regression: high variance

Regression 1.




Regression: high variance

Regression 2.



Regression: high variance

Average of regression 1 and regression 2 gives better prediction.



Majority voting of classifiers

- Consider M classifiers $f_1(x), \dots, f_M(x)$, performing binary classification.
- Let probability of mistake be constant $p \in (0, \frac{1}{2})$:
 $p(f_m(x) \neq y) = p \forall m$
- Suppose all models make mistakes or correct guesses independently of each other.
- Let $G(x)$ be majority voting combiner.
- Then $p(G(x) \neq y) \rightarrow 0$ as $m \rightarrow \infty$

Convex loss

Convex loss promotes the usage of averaged prediction instead of individual ones.

- Take convex loss $\mathcal{L}(\hat{y} - y)$, such as absolute or square.
- Take $f_1(x), \dots, f_M(x)$ with weights w_1, \dots, w_M .
- For any fixed x consider 2 prediction strategies:
 - 1 sample $m \sim \text{Categorical}(\alpha_1, \dots, \alpha_M)$, $\hat{y}(x) = f_m(x)$.
 - 2 $\hat{y}(x) = \sum_{m=1}^M w_m f_m(x)$
- Second strategy is better than first¹, averaged over different sample outcomes m .

¹Prove that.

Ambiguity decomposition

Ambiguity decomposition:

consider predicting fixed (x, y) with ensemble for

$F(x) = \sum_{m=1}^M w_m f_m(x)$, $w_m \geq 0$, $\sum_m w_m = 1$. Then

$$\underbrace{(F(x) - y)^2}_{\text{ensemble error}} = \underbrace{\sum_m w_m (f_m(x) - y)^2}_{\text{base learner error}} - \underbrace{\sum_m w_m (f_m(x) - F(x))^2}_{\text{ambiguity}}$$

Ensemble is accurate when:

- $f_m(x)$ are accurate
- and/or there is huge disagreement in base learners predictions.

Proof of ambiguity decomposition

Proof:

$$\begin{aligned} \sum_m w_m (f_m(x) - F(x))^2 &= \sum_m w_m (f_m(x) - y + y - F(x))^2 \\ &= \sum_m w_m (f_m(x) - y)^2 + \sum_m w_m (y - F(x))^2 + 2 \sum_m w_m (f_m(x) - y) (y - F(x)) \\ &= \sum_m w_m (f_m(x) - y)^2 + (F(x) - y)^2 + 2 (y - F(x)) \sum_m w_m (f_m(x) - y) \\ &= \sum_m w_m (f_m(x) - y)^2 + (F(x) - y)^2 + 2 (y - F(x)) (F(x) - y) \\ &= \sum_m w_m (f_m(x) - y)^2 + (F(x) - y)^2 - 2 (F(x) - y)^2 \\ &= \sum_m w_m (f_m(x) - y)^2 - (F(x) - y)^2 \end{aligned}$$

Data promotes ensembles

Data may promote the use of ensembles when it is divided into separate groups with different regularities.

- Flat price prediction:
 - purpose-for living: model depending on comfort, living tastes, etc.
 - purpose-for investment: another model depending on exchange rates, interest rates, stock growth, etc.
- Face detection on images:
 - one model detects face with frontal view
 - another model detects face with profile view
- Person identification using diverse information:
 - by voice, by face, by behaviour patterns, etc.

Table of Contents

- 1 Bias-variance decomposition
- 2 Ensemble learning
- 3 Fixed integration schemes for classification**
- 4 Stacking
- 5 Train set sampling ensemble methods

Classifiers output **labels**

- **Binary classification:** output $+1 \Leftrightarrow$
 - all classifiers predict $+1$ (AND rule)
 - at least one classifier predicts $+1$ (OR rule)
 - at least k classifiers predict $+1$ (k-out-of-N)
- **Multiclass classification:**
 - predict most popular class (majority vote)
- **Extension - weighted account for classifiers:**
 - weighted majority vote
 - weighted k-out-of-N

Classifiers output **scores**

- Let $g_y^m(x)$ be score of class y by model m .
- Problem: scores are incomparable across models.
- Solution:
 - 1 define ranking score: $s_y^m(x) = \sum_{c \neq y} \mathbb{I}[g_y^m(x) > g_c^m(x)]$
 - 2 since $s_y^m(x)$ are comparable, assign

$$\hat{y}(x) = \arg \max_y \sum_{m=1}^M s_y^m(x)$$

Allows weighted account of classifiers.

Classifiers output **probabilities**

- Let $p_y^m(x)$ be probability of class y by classifier m .
- Possible final predictions:

$$p_y(x) = \frac{1}{M} \sum_{m=1}^M p_y^m(x)$$

$$p_y(x) = \text{median}_m p_y^m(x)$$

Allows weighted account of classifiers.

Table of Contents

- 1 Bias-variance decomposition
- 2 Ensemble learning
- 3 Fixed integration schemes for classification
- 4 Stacking**
- 5 Train set sampling ensemble methods

Stacking algorithm

Input:

- original training set $T = \{(x_n, y_n)\}_{n=1}^N$
- base learners $f_1(x), \dots, f_M(x)$ and $G(\cdot)$.

Stacking algorithm

Input:

- original training set $T = \{(x_n, y_n)\}_{n=1}^N$
- base learners $f_1(x), \dots, f_M(x)$ and $G(\cdot)$.
- ① Set generated training set $T' = \{\}$
- ② Split training set into K folds: T_1, T_2, \dots, T_K .

Stacking algorithm

Input:

- original training set $T = \{(x_n, y_n)\}_{n=1}^N$
 - base learners $f_1(x), \dots, f_M(x)$ and $G(\cdot)$.
- 1 Set generated training set $T' = \{\}$
 - 2 Split training set into K folds: T_1, T_2, \dots, T_K .
 - 3 for k in $1, 2, \dots, K$:
 - train $f_1(x), \dots, f_M(x)$ on $T \setminus T_k$
 - for (x, y) in T_k :
 - augment T' with sample $([f_1(x), \dots, f_M(x)], y)$

Stacking algorithm

Input:

- original training set $T = \{(x_n, y_n)\}_{n=1}^N$
- base learners $f_1(x), \dots, f_M(x)$ and $G(\cdot)$.
- ① Set generated training set $T' = \{\}$
- ② Split training set into K folds: T_1, T_2, \dots, T_K .
- ③ for k in $1, 2, \dots, K$:
 - train $f_1(x), \dots, f_M(x)$ on $T \setminus T_k$
 - for (x, y) in T_k :
 - augment T' with sample $([f_1(x), \dots, f_M(x)], y)$
- ④ Train $G(\cdot)$ on T' .

Stacking algorithm

Input:

- original training set $T = \{(x_n, y_n)\}_{n=1}^N$
- base learners $f_1(x), \dots, f_M(x)$ and $G(\cdot)$.
- ❶ Set generated training set $T' = \{\}$
- ❷ Split training set into K folds: T_1, T_2, \dots, T_K .
- ❸ for k in $1, 2, \dots, K$:
 - train $f_1(x), \dots, f_M(x)$ on $T \setminus T_k$
 - for (x, y) in T_k :
 - augment T' with sample $([f_1(x), \dots, f_M(x)], y)$
- ❹ Train $G(\cdot)$ on T' .
- ❺ Retrain $f_1(x), \dots, f_M(x)$ on T .

Output: ensemble $G(f_1(x), \dots, f_M(x))$.

Comments

- Training $f_1(x), \dots, f_M(x), G(\cdot)$ on the same data causes overfitting.
- Besides $f_1(x), \dots, f_M(x)$ $G(\cdot)$ may also depend on
 - original features x
 - internal representations inside f_m such as class scores, probabilities.

Linear stacking (blending)

- Linear stacking:

$$f(x) = \sum_{m=1}^M w_m f_m(x)$$

$$\left(\sum_{m=1}^M w_m f_m(x_n) - y_n \right)^2 \rightarrow \min_{\mathbf{w}}$$

- $f_1(x), \dots, f_M(x)$ are correlated (predict the same y) \Rightarrow estimate unstable.
- For more robust estimate solve:

$$\begin{cases} \left(\sum_{m=1}^M w_m f_m(x_n) - y_n \right)^2 + \lambda \sum_{m=1}^M \left(w_m - \frac{1}{M} \right)^2 \rightarrow \min_{\mathbf{w}} \\ w_1 \geq 0, \dots, w_M \geq 0 \end{cases}$$

Table of Contents

- 1 Bias-variance decomposition
- 2 Ensemble learning
- 3 Fixed integration schemes for classification
- 4 Stacking
- 5 Train set sampling ensemble methods

Bagging & random subspaces

When model overfits to particular training set T , it is useful to generate many training sets T_1, \dots, T_M , estimate model on each of them and average.

- **Bagging:**
 - random selection of samples (with replacement)²³
- **Random subspace method:**
 - random selection of features (without replacement)
- May apply both methods jointly.

²what is the probability that observation will not belong to bootstrap sample?

³what is the limit of this probability with $N \rightarrow \infty$?

Bagged trees

In CART trees we solve

$$\hat{f}, \hat{h} = \arg \min_{f, h \in S(t)} \Delta I(t)$$

$S(t)$ for standard decision trees:

```
S = {}  
for each f in {1, ..., D}  
  for each h in unique {x_n^f}_{n: x_n \in t}  
    S := S \cup (f, h)
```

Bagged decision trees - bagging applied to standard decision trees.

Random forest & extra random trees

- Random forest & extra random trees are bagged decision trees with restricted search through (f, h) , controlled by $\alpha \in (0, 1]$.
 - restricted search \Rightarrow higher bias, smaller variance.

$S(t)$ for random forest:

```

 $S = \{\}, K = \alpha D$ 
sample  $d_1, \dots, d_K$  randomly from  $\{1, \dots, D\}$  without replacement.
for each  $f$  in  $d_1, \dots, d_K$ 
  for each  $h$  in  $\text{unique} \{x_n^f\}_{n: x_n \in t}$ 
     $S := S \cup (f, h)$ 
  
```

$S(t)$ for extra random trees:

```

 $S = \{\}, K = \alpha D$ 
sample  $d_1, \dots, d_K$  randomly from  $\{1, \dots, D\}$  without replacement.
for each  $f$  in  $d_1, \dots, d_K$ 
  sample  $h$  randomly from  $\text{unique} \{x_n^f\}_{n: x_n \in t}$ 
   $S := S \cup (f, h)$ 
  
```


Out-of-bag estimate

Out-of-bag estimate - estimate of expected loss by bagged algorithms.

- from above (pessimistic)
- without need for separate validation set

$$OOB = \frac{1}{N} \sum_{n=1}^N \mathcal{L} \left(\frac{1}{|I_n|} \sum_{m \in I_n} f_m(x_n), y_n \right)$$


For each (x_n, y_n) we make prediction using only those models $I_n = \{m : (x_n, y_n) \notin T_m\}$ for which (x_n, y_n) is new.

Comments

- Bagged decision trees, RF, ERT:
 - have straightforward parallel implementation
 - but trees are not targeted to correct mistakes of each other
- Trees in RF, ERT may be built on the same training set T
 - due to stochastic $S(t)$ they will be different anyway

Comments

- Let $M = \#$ of base learners.
- ERT trains faster than RF, but on average requires higher M .
- Typical dependency between loss of bagging/RF/ERT depending on M :



- We average variability of tree to training set, what is more efficient for higher M .
- To find optimal hyperparameters set small M , find other parameters, then set high M back.

Conclusion

- Bias-variance decomposition gives 2 sources for poor accuracy:
 - bias: for underfitted models
 - variance: for overfitted models
- Stacking with complex aggregating model decreases bias.
 - may add variance
- Stacking with simple aggregating model (averaging, majority vote) decreases variance.