

Boosting

Victor Kitov

v.v.kitov@yandex.ru



Linear ensembles

Linear ensemble:

$$F_M(x) = f_0(x) + c_1 f_1(x) + \dots + c_M f_M(x)$$

Regression: $\hat{y}(x) = F_M(x)$

Binary classification: $\text{score}(y|x) = F_M(x)$, $\hat{y}(x) = \text{sign } F_M(x)$

- Notation: $f_1(x), \dots, f_M(x)$ are called *base learners*, *weak learners*, *base models*.
- Too expensive to optimize $f_0(x), f_1(x), \dots, f_M(x)$ and c_1, \dots, c_M jointly for large M .
- Idea: optimize $f_0(x)$ and then each pair $(f_m(x), c_m)$ greedily.

Forward stagewise additive modeling (FSAM)

Input:

- training dataset (x_n, y_n) , $n = 1, 2, \dots, N$
- loss function $\mathcal{L}(f, y)$
- parametric form of base learner $f(x|\gamma)$ (parametrized by γ)
- the number of base learners M .

Output: approximation function $F_M(x) = f_0(x) + \sum_{m=1}^M c_m f_m(x)$

Forward stagewise additive modeling (FSAM)

- ① Fit initial approximation $f_0(x) = \arg \min_f \sum_{n=1}^N \mathcal{L}(f(x_n), y_n)$
- ② For $m = 1, 2, \dots, M$:
 - find next best classifier

$$(c_m, f_m) = \arg \min_{f, c} \sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) + cf(x_n), y_n)$$

- reevaluate ensemble

$$F_m(x) = F_{m-1}(x) + c_m f_m(x)$$

Comments

- M should be determined by performance on validation set.
 - may overfit!
- Each step should be coarse to leave room for future base learners improvement:
 - initial approximation may be zero or constant
 - optimization can be coarse (just few steps)
 - base learner should be simple
 - such as trees of depth=1,2,3.
- For some loss functions (see Adaboost) we can solve minimization explicitly.
- For general loss functions gradient boosting should be used.

Table of Contents

1 Adaboost

2 Gradient boosting

3 Extensions

Adaboost (discrete version)

Assumptions:

- binary classification task $y \in \{+1, -1\}$
- $f_m(x) \in \{+1, -1\}$
- classification is performed with
$$\hat{y} = \text{sign}\{f_0(x) + c_1 f_1(x) + \dots + c_M f_M(x)\}$$
- optimized loss is $\mathcal{L}(F(x), y) = e^{-yF(x)}$

Optimization in FSAM can be solved explicitly!

Adaboost (discrete version): algorithm

Input:

- training dataset (x_n, y_n) , $n = 1, 2, \dots, N$
- number of additive weak classifiers M
- a family of weak classifiers $h(x) \in \{+1, -1\}$
 - should be trainable on weighted datasets.

Output: composite classifier $F_M(x) = \text{sign} \left(\sum_{m=1}^M c_m f_m(x) \right)$

Adaboost (discrete version): algorithm

- 1 Initialize observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.
- 2 for $m = 1, 2, \dots, M$:

- 1 fit $f_m(x)$ to training data using weights w_i
- 2 compute weighted misclassification rate:

$$E_m = \frac{\sum_{i=1}^N w_i \mathbb{I}[f_m(x) \neq y_i]}{\sum_{i=1}^N w_i}$$

- 3 compute weighting factor $c_m = \frac{1}{2} \ln((1 - E_m)/E_m)$
- 4 if $E_m > 0.5$ or $E_m = 0$: terminate procedure.
- 5 increase all weights, where misclassification with $f_m(x)$ was made:

$$w_i \leftarrow w_i \frac{1 - E_m}{E_m}, \quad i \in \{i : f_m(x_i) \neq y_i\}$$

- 3 Return classifier $F_M(x) = \text{sign} \left(\sum_{m=1}^M c_m f_m(x) \right)$

Table of Contents

1 Adaboost

2 Gradient boosting

3 Extensions

Motivation

- Problem: For general loss function L FSAM cannot be solved explicitly
- Analogy with function minimization: when we can't find optimum explicitly we use numerical methods
- Gradient boosting: numerical method for iterative loss minimization

Gradient descent algorithm

$$L(w) \rightarrow \min_w, \quad g(w) = \nabla_w L(w), \quad w \in \mathbb{R}^N$$

Gradient descend algorithm:

- given w move in the direction of steepest descent, given by $\Delta w := -g(w)$, $g(w) = \nabla_w L(w)$

Gradient descent algorithm

$$L(w) \rightarrow \min_w, \quad g(w) = \nabla_w L(w), \quad w \in \mathbb{R}^N$$

INPUT:

step size ε

number of iterations M

ALGORITHM:

initialize w

for $m = 1, 2, \dots, M$:

$\Delta w = -g(w)$

$w = w + \varepsilon \Delta w$

Modified gradient descent algorithm

$$L(w) \rightarrow \min_w, \quad g(w) = \nabla_w L(w), \quad w \in \mathbb{R}^N$$

INPUT:

number of iterations M

ALGORITHM:

initialize w

for $m = 1, 2, \dots, M$:

$$\Delta w = -g(w)$$

$$c^* = \arg \min_{c > 0} L(w + c\Delta w)$$

$$w = w + c^* \Delta w$$

Gradient boosting

- Now consider

$$L(f(x_1), \dots, f(x_N)) = \sum_{n=1}^N \mathcal{L}(f(x_n), y_n) \rightarrow \min_{f(\cdot)}$$

- Gradient descent performs pointwise optimization, but we need generalization, so we optimize in space of functions.
- Gradient boosting = modified gradient descent in function space:
 - find $z_n = -g(x_n)$, where $g(x_n) = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \big|_{r=f^{m-1}(x_n)}$
 - fit base learner $f_m(x)$ to $\{(x_n, z_n)\}_{n=1}^N$

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- 2 For each step $m = 1, 2, \dots, M$:

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- ② For each step $m = 1, 2, \dots, M$:
 - ① calculate targets $z_n := -g_n$ $\left(g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)$

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ❶ Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- ❷ For each step $m = 1, 2, \dots, M$:
 - ❶ calculate targets $z_n := -g_n \quad \left(g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)$
 - ❷ fit $f_m(\cdot)$ to $\{(x_n, z_n)\}_{n=1}^N$, for example by solving

$$\sum_{n=1}^N (f_m(x_n) - z_n)^2 \rightarrow \min_{f_m}$$

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- ② For each step $m = 1, 2, \dots, M$:
 - ① calculate targets $z_n := -g_n \quad \left(g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)$
 - ② fit $f_m(\cdot)$ to $\{(x_n, z_n)\}_{n=1}^N$, for example by solving

$$\sum_{n=1}^N (f_m(x_n) - z_n)^2 \rightarrow \min_{f_m}$$

- ③ solve univariate optimization problem:

$$\sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) + c_m f_m(x_n), y_n) \rightarrow \min_{c_m > 0}$$

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- ② For each step $m = 1, 2, \dots, M$:

- ① calculate targets $z_n := -g_n \quad \left(g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)$
- ② fit $f_m(\cdot)$ to $\{(x_n, z_n)\}_{n=1}^N$, for example by solving

$$\sum_{n=1}^N (f_m(x_n) - z_n)^2 \rightarrow \min_{f_m}$$

- ③ solve univariate optimization problem:

$$\sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) + c_m f_m(x_n), y_n) \rightarrow \min_{c_m > 0}$$

- ④ set $F_m(x) = F_{m-1}(x) + c_m f_m(x)$

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- ② For each step $m = 1, 2, \dots, M$:
 - ① calculate targets $z_n := -g_n \quad \left(g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)$
 - ② fit $f_m(\cdot)$ to $\{(x_n, z_n)\}_{n=1}^N$, for example by solving

$$\sum_{n=1}^N (f_m(x_n) - z_n)^2 \rightarrow \min_{f_m}$$

- ③ solve univariate optimization problem:

$$\sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) + c_m f_m(x_n), y_n) \rightarrow \min_{c_m > 0}$$

- ④ set $F_m(x) = F_{m-1}(x) + c_m f_m(x)$

Output: approximation function $F_M(x) = f_0(x) + \sum_{m=1}^M c_m f_m(x)$

Gradient boosting: examples

In gradient boosting

$$\sum_{n=1}^N \left(f_m(x_n) - \left(-\frac{\partial \mathcal{L}(r, y)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right) \right)^2 \rightarrow \min_{f_m}$$

Sample cases:

- $\mathcal{L} = \frac{1}{2} (r - y)^2$
- $\mathcal{L} = [-ry]_+$

Table of Contents

1 Adaboost

2 Gradient boosting

3 Extensions

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- 2 For each step $m = 1, 2, \dots, M$:

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- 2 For each step $m = 1, 2, \dots, M$:

- 1 calculate targets $z_n := -g_n$ $\left(g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)$

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- 2 For each step $m = 1, 2, \dots, M$:

- 1 calculate targets $z_n := -g_n$ $\left(g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)$
- 2 fit regression tree $f_m(\cdot)$ on $\{(x_n, z_n)\}_{n=1}^N$ with some loss function, get leaf regions $\{R_j^m\}_{j=1}^{J_m}$.

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- ② For each step $m = 1, 2, \dots, M$:

- ① calculate targets $z_n := -g_n \quad \left(g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)$
- ② fit regression tree $f_m(\cdot)$ on $\{(x_n, z_n)\}_{n=1}^N$ with some loss function, get leaf regions $\{R_j^m\}_{j=1}^{J_m}$.
- ③ for each terminal region R_j^m , $j = 1, 2, \dots, J_m$ solve univariate optimization problem:

$$\gamma_j^m = \arg \min_{\gamma} \sum_{x_i \in R_j^m} \mathcal{L}(F_{m-1}(x_i) + \gamma, y_i)$$

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- ② For each step $m = 1, 2, \dots, M$:

- ① calculate targets $z_n := -g_n \quad \left(g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)$
- ② fit regression tree $f_m(\cdot)$ on $\{(x_n, z_n)\}_{n=1}^N$ with some loss function, get leaf regions $\{R_j^m\}_{j=1}^{J_m}$.
- ③ for each terminal region R_j^m , $j = 1, 2, \dots, J_m$ solve univariate optimization problem:

$$\gamma_j^m = \arg \min_{\gamma} \sum_{x_i \in R_j^m} \mathcal{L}(F_{m-1}(x_i) + \gamma, y_i)$$

- ④ update $F_m(x) = F_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_j^m \mathbb{I}[x \in R_j^m]$

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- ② For each step $m = 1, 2, \dots, M$:

- ① calculate targets $z_n := -g_n \quad \left(g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)$
- ② fit regression tree $f_m(\cdot)$ on $\{(x_n, z_n)\}_{n=1}^N$ with some loss function, get leaf regions $\{R_j^m\}_{j=1}^{J_m}$.
- ③ for each terminal region R_j^m , $j = 1, 2, \dots, J_m$ solve univariate optimization problem:

$$\gamma_j^m = \arg \min_{\gamma} \sum_{x_i \in R_j^m} \mathcal{L}(F_{m-1}(x_i) + \gamma, y_i)$$

- ④ update $F_m(x) = F_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_j^m \mathbb{I}[x \in R_j^m]$

Output: approximation function $F_M(x)$

Modification of boosting for trees

- Compared to first method of gradient boosting, boosting of regression trees finds additive coefficients individually for each terminal region R_j^m , not globally for the whole classifier $f_m(x)$.
- This is done to increase accuracy: forward stagewise algorithm cannot be applied to find R_j^m , but it can be applied to find γ_j^m , because second task is solvable for arbitrary L .
- Max leaves J
 - interaction between no more than $J - 1$ terms
 - usually $4 \leq J \leq 8$
- M controls underfitting-overfitting tradeoff and selected using validation set

Shrinkage & subsampling

- Shrinkage of general GB, step (d):

$$F_m(x) = F_{m-1}(x) + \alpha c_m f_m(x)$$

- Comments:
 - $\alpha \in (0, 1]$
 - $\alpha \downarrow \implies M \uparrow (\alpha M \approx \text{const})$
- Subsampling
 - increases speed of fitting
 - may increase accuracy

Linear loss function approximation

Consider sample (x, y) and $g(x) = \left. \frac{\partial \mathcal{L}(r, y)}{\partial r} \right|_{r=F(x)}$

$$\mathcal{L}(F(x) + f(x), y) \approx \mathcal{L}(F(x), y) + g(x)f(x)$$

$\Rightarrow f(x)$ should be fitted to $-g(x)$.

Quadratic loss function approximation

$$\begin{aligned}\text{Define } g(x) &= \left. \frac{\partial \mathcal{L}(r, y)}{\partial r} \right|_{r=F(x)}, \quad h(x) = \left. \frac{\partial^2 \mathcal{L}(r, y)}{\partial r^2} \right|_{r=F(x)} \\ \mathcal{L}(F(x) + f(x), y) &\approx \\ \mathcal{L}(F(x), y) + g(x)f(x) + \frac{1}{2}h(x)(f(x))^2 &= \\ = \frac{1}{2}h(x) \left(f(x) + \frac{g(x)}{h(x)} \right)^2 + \text{const}(f(x))\end{aligned}$$

So $f(x)$ should be fitted to $-g(x)/h(x)$ with weight $h(x)$.

- $h(x) \geq 0$ around local minimum.

Case $y \in \{1, 2, \dots, C\}$

One-vs-all, one-vs-one, error-correcting-codes.

Case $y \in \{1, 2, \dots, C\}$

One-vs-all, one-vs-one, error-correcting-codes.

Alternatively can optimize $\mathcal{L}(F(x), y)$ for $F(x) \in \mathbb{R}^C$

- $F(x) = \{p(y = c|x)\}_{c=1}^C$, y - one-hot encoded true class
- $\mathcal{L}(F(x), y) = F(x)^T y = p(y = \text{correct class}|x)$
- $z_n = -\frac{\partial \mathcal{L}(r, y)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \in \mathbb{R}^C$
- $\sum_{n=1}^N (f_m(x_n) - z_n)^2 \rightarrow \min_{f_m}$ yields vector C -dim. regression.
- may use quadratic approximation
 - for efficient inverting of $\left(\frac{\partial^2 \mathcal{L}(r, y)}{\partial r^2} \Big|_{r=F(x)} \right)$ may use diagonal approximation.

xgBoost

- One of the most popular algorithms on kaggle.
- Uses decision trees as base learners:
 - $f_m \in \{f(x) = w_{q(x)}\}$,
 - T total number of leaves.
 - $q(x)$ maps $x \in \mathbb{R}^D$ to leaf number
 - $w \in \mathbb{R}^T$ predictions for leaves.

xgBoost

- Loss - 2nd order approximation with **with regularization**:

$$\begin{aligned}\mathcal{L}(f_m) &= \sum_{n=1}^N \mathcal{L}(F^{(m-1)}(x_n), y_n) \\ &\approx \sum_{n=1}^N \left[\mathcal{L}(F^{(m-1)}(x_n), y_n) + g_n f_m(x_n) + \frac{1}{2} h_n f_m^2(x_n) \right] \\ &\quad + \gamma T + \frac{1}{2} \lambda \sum_{t=1}^T w_t^2\end{aligned}$$

- Tree impurity function matches original loss $\mathcal{L}(\cdot, \cdot)$.
- Efficiency optimization:
 - feature values may be discretized for speed
 - parallelization over multiple CPU cores and with GPU

Types of boosting

- Loss function \mathcal{L} :
 - $\mathcal{L}(|f(x) - y|)$ - regression
 - $F(y \cdot \text{score}(y = +1|x))$ - binary classification
 - $\mathcal{L}(F(x), y)$ for $F(x), y \in \mathbb{R}^C$ - multiclass classification
- Optimization
 - analytical (Adaboost)
 - gradient based
 - based on quadratic approximation
- Base learners
 - continious
 - discrete
- Classification
 - binary
 - multiclass
- Extensions: shrinkage, subsampling