

# Recurrent neural nets.

Victor Kitov

[v.v.kitov@yandex.ru](mailto:v.v.kitov@yandex.ru)



School of Data Analysis.

# Intro

- Sequences
  - words: sequences of symbols
  - sentences: sequences of words
  - documents: sequences of words
- Need fixed vector representation for prediction!
- Bag-of-words allows to do that:
  - one-hot encoding: indicator, TF, TF-IDF models
  - embeddings: get average embedding for sentence/document.

## Problem of bag-of-words approach

- Problem: bag-of-words completely ignores word order.
  - information loss!
- Recurrent neural nets account for positions of all elements in sequence!
  - output fixed size sequence representation
  - this feature representation is feature extraction for later model.
    - e.g. MLP.

# Recurrent neural net (RNN)

- Consider input sequence  $\mathbf{x}_{1:n} := \mathbf{x}_1, \dots, \mathbf{x}_n$ ,  $\mathbf{x}_i \in \mathbb{R}^{d_{in}}$ .
- RNN outputs single vector  $\hat{\mathbf{y}}_n \in \mathbb{R}^{d_{out}}$ :

$$\hat{\mathbf{y}}_n = RNN(\mathbf{x}_{1:n})$$

- This implicitly defines  $RNN^*$  with sequential output:

$$\hat{\mathbf{y}}_{1:n} = RNN^*(\mathbf{x}_{1:n})$$

$$\hat{\mathbf{y}}_i = RNN(\mathbf{x}_{1:i})$$


- Comments:
  - RNN shrinks history  $\mathbf{x}_{1:n}$  to fixed size vector  $\mathbf{y}_n$ .
  - No Markov assumption: all info is aggregated!

# Technical details of RNN

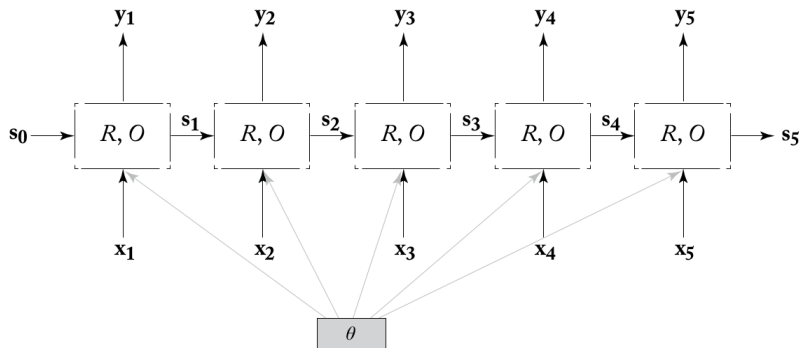
$$\begin{aligned}
 RNN^*(\mathbf{x}_{1:n}, \mathbf{s}_0) &= \mathbf{y}_{1:n} \\
 \hat{\mathbf{y}}_i &= O(\mathbf{s}_i) \\
 \mathbf{s}_i &= R(\mathbf{s}_{i-1}, \mathbf{x}_i)
 \end{aligned}$$

$$\mathbf{x}_i \in \mathbb{R}^{d_{in}}, \mathbf{y}_i \in \mathbb{R}^{d_{out}}, \mathbf{s}_i \in \mathbb{R}^{d_{state}}$$

Typical usage:  $O(\mathbf{s}) \equiv \mathbf{s}$ ,  $d_{state} = d_{out}$ ,  $\mathbf{s}_0 = \mathbf{0}$ .



## Unrolled RNN



$$\begin{aligned}
 s_4 &= R(s_3, x_4) = R(R(s_2, x_3), x_4) \\
 &= R(R(R(s_1, x_2), x_3), x_4) = R(R(R(R(s_0, x_1), x_2), x_3), x_4)
 \end{aligned}$$

# Training

Training: unroll RNN and use parameter sharing.

- called **backpropagation through time** (BPTT)
- Variant: unroll RNN for all non-intersecting subsequences of given sequence of given length.

```

init  $s_0$ 

for  $i$  in  $0, 1, \dots, n/k - 1$ :
     $\hat{y}_{ki+1:ki+k} = RNN^*(x_{ki+1:ki+k}, s_{ki})$ 
    calculate loss  $\sum_{j=ki+1}^{ki+k} L(\hat{y}_j, y_j)$ 
    backpropagate gradients, update weights
  
```

Mostly used for simple RNN, as gated RNN remember events long ago.

## Common use-cases of RNN

- **Acceptor:** output prediction in  $\hat{\mathbf{y}}_n$ .
  - e.g. read sentence and output its polarity probabilities.
- **Encoder:** encode input sequence representation as  $\hat{\mathbf{y}}_n$ , e.g.:
  - machine translation: translation done with another "decoding" RNN  
decode starting from  $\mathbf{s}_0 = \hat{\mathbf{y}}_n$ .
  - summarization: for each sentence classify whether to include it into summary or not.  
besides sentence features use  $\hat{\mathbf{y}}_n$  as document summary feature in classification.
- **Transducer:** tag sequence  $x_1, \dots, x_n$  with RNN outputs  $y_1, \dots, y_n$ .  
Loss function:


$$\mathcal{L}(\hat{\mathbf{y}}_{1:n}, \mathbf{y}_{1:n}) = \sum_{i=1}^n L(\hat{\mathbf{y}}_i, \mathbf{y}_i)$$

- e.g. POS tagging, language modelling.



# NN & RNN architectures

NN & RNN architectures:



Examples, where these architectures arise:

- **one to one:** classical classification, image classification.
- **one to many:** image captioning, story generation based on topic.
- **many to one:** text classification, sentiment analysis.
- **many to many:** machine translation, summarization.
- **syncd many to many:** POS tagging, activity detection on video.

# Table of Contents

- 1 RNN extensions
- 2 Concrete RNN Architectures
- 3 LSTM model

# Bidirectional RNN

- Bidirectional RNN consists of 2 RNNs
  - forward RNN ( $R^f, O^f$ ) with state  $\mathbf{s}_i^f, i = \overline{1, n}$
  - backward RNN ( $R^b, O^b$ ) with state  $\mathbf{s}_i^b, i = \overline{1, n}$
- Forward RNN goes in forward direction  $\mathbf{x}_1, \mathbf{x}_2 \dots \mathbf{x}_n$ .
- Backward RNN goes in backward direction  $\mathbf{x}_n, \mathbf{x}_{n-1} \dots \mathbf{x}_1$ .
- At each moment  $i$  we have 2 states:
  - 1  $\mathbf{s}_i^f = F_1(\mathbf{x}_1, \mathbf{x}_2 \dots \mathbf{x}_i)$
  - 2  $\mathbf{s}_i^b = F_2(\mathbf{x}_i, \mathbf{x}_{i+1}, \dots \mathbf{x}_n)$

# Bidirectional RNN

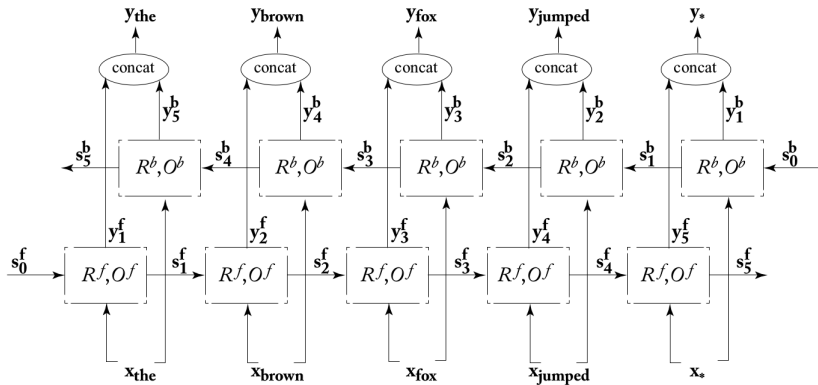
- So we can output

$$biRNN(\mathbf{x}_{1:n}, i) = \hat{\mathbf{y}}_i = [\hat{\mathbf{y}}_i^f; \hat{\mathbf{y}}_i^b] = [RNN^f(\mathbf{x}_{1:i}); RNN^b(\mathbf{x}_{n:i})]$$

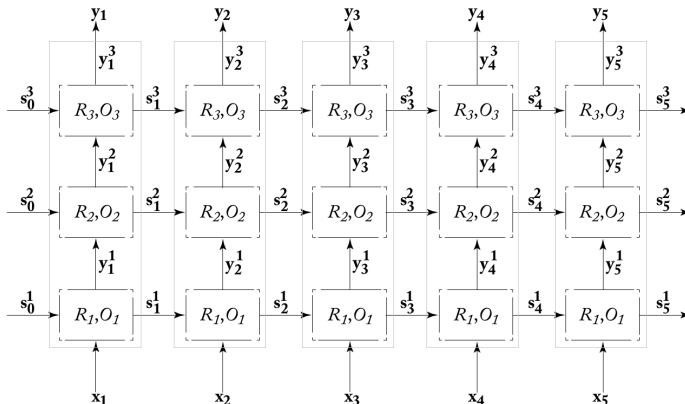
$$biRNN^*(\mathbf{x}_{1:n}) = \mathbf{y}_{1:n} = [biRNN(\mathbf{x}_{1:n}, 1); \dots; biRNN(\mathbf{x}_{1:n}, n)]$$

- encoding takes into account past and future!
- biRNN is very effective for tagging sequences (e.g. POS tagging).

## biRNN illustration



# Stacked RNN



- Output of previous layer RNN is input to next layer.
- Empirically stacked RNNs work better than single layer RNNs.
- biRNNs can also be stacked.

# Table of Contents

- 1 RNN extensions
- 2 Concrete RNN Architectures**
- 3 LSTM model

## Bag-of-words RNN

Bag-of-words RNN:

$$\mathbf{s}_i = \mathbf{s}_{i-1} + \mathbf{x}_i$$

$$\mathbf{y}_i = \mathbf{s}_i$$

- $\mathbf{x}_i$ : input vector
- $\mathbf{s}_i$ : hidden layer
- $\mathbf{y}_i$ : output vector

Order of words does not matter, not very informative.



## Simple RNN (S-RNN)

Simple RNN (S-RNN)<sup>1</sup>:

$$\begin{aligned}\mathbf{s}_i &= g_s(W_s \mathbf{s}_{i-1} + V_s \mathbf{x}_i + \mathbf{b}_s) \\ \mathbf{y}_i &= g_y(W_y \mathbf{s}_i + \mathbf{b}_y)\end{aligned}$$

- $\mathbf{x}_i$ : input vector
- $\mathbf{s}_i$ : hidden layer
- $\mathbf{y}_i$ : output vector
- $W_s, V_s, W_y$ : parameter matrices
- $\mathbf{b}_s, \mathbf{b}_y$ : parameter vectors
- $g_s(\cdot), g_y(\cdot)$ : activation functions

---


<sup>1</sup>also called Elman network

## Properties of S-RNN

- S-RNN is sensitive to the order of the inputs.
- Due to recurrent multiplications by  $W_s$  is subject to:
  - exploding gradient problem
    - solved by gradient clipping
  - vanishing gradient problem
    - solved by gated models and memory models

# Exploding gradient problem

Exploding gradient problem:



# Exploding gradient problem

Solutions:

- add regularization
- gradient clipping: clip norm of gradient by threshold.
  - if  $\|\nabla_{\theta} L(\hat{\mathbf{y}}_i, \mathbf{y}_i)\| < t$

$$\theta \rightarrow \theta - \varepsilon \nabla_{\theta} L(\hat{\mathbf{y}}_i, \mathbf{y}_i)$$

- else

$$\theta \rightarrow \theta - \varepsilon \frac{t}{\|\nabla_{\theta} L(\hat{\mathbf{y}}_i, \mathbf{y}_i)\|} \nabla_{\theta} L(\hat{\mathbf{y}}_i, \mathbf{y}_i)$$

# Vanishing gradients

- Repetitive multiplication of state by the same matrix  $W_s$  and saturating non-linearities also cause net to forget past quickly due to vanishing gradients.
- Ways to combat this:
  - 1 Initialize  $W_s = I, b_s = 0, g_s = ReLu$ .
    - so initially network sums information
    - will change behavior after training if needed
  - 2 Better solution: use LSTM model.

# Table of Contents

- 1 RNN extensions
- 2 Concrete RNN Architectures
- 3 LSTM model

# Gates

- Consider  $n$  dimensional vectors:
  - old state  $\mathbf{s}$ , update  $\mathbf{x}$  and new state  $\mathbf{s}'$ .
- Gate  $g \in \{0, 1\} \in \mathbb{R}^n$  controls state positions where change is applied.
- Example ( $\odot$  defines point-wise multiplication):

$$\begin{array}{c} \begin{bmatrix} 8 \\ 11 \\ 3 \\ 7 \\ 5 \\ 15 \end{bmatrix} \\ \mathbf{s}' \end{array} \leftarrow \begin{array}{c} \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 1 \end{bmatrix} \\ \mathbf{g} \end{array} \odot \begin{array}{c} \begin{bmatrix} 10 \\ 11 \\ 12 \\ 13 \\ 14 \\ 15 \end{bmatrix} \\ \mathbf{x} \end{array} + \begin{array}{c} \begin{bmatrix} 1 \\ 0 \\ 1 \\ 1 \\ 1 \\ 0 \end{bmatrix} \\ \mathbf{(1-g)} \end{array} \odot \begin{array}{c} \begin{bmatrix} 8 \\ 9 \\ 3 \\ 7 \\ 5 \\ 8 \end{bmatrix} \\ \mathbf{s} \end{array}$$

- Problems:
  - gates need to be learned
  - piece-wise constant gates cannot be optimized.
- Solution: use sigmoid gate  $\mathbf{g} = \sigma(f(\mathbf{x}, \mathbf{s}, \theta))$ 
  - $\theta$ : learned parameters
  - $f$ : any differentiable function

# Long short-term memory (LSTM) model

$$\mathbf{f}_t = \sigma(W_f \mathbf{x}_t + U_f \mathbf{h}_{t-1} + \mathbf{b}_f)$$

forget gate

$$\mathbf{i}_t = \sigma(W_i \mathbf{x}_t + U_i \mathbf{h}_{t-1} + \mathbf{b}_i)$$

input gate

$$\mathbf{o}_t = \sigma(W_o \mathbf{x}_t + U_o \mathbf{h}_{t-1} + \mathbf{b}_o)$$

output gate

$$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \tanh(W_c \mathbf{x}_t + U_c \mathbf{h}_{t-1} + \mathbf{b}_c)$$

inner state

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c}_t)$$

observed output


$\mathbf{x}_t$ -input, parameters:

- matrices:  $W_f, U_f, W_i, U_i, W_o, U_o, W_c, U_c$
- vectors:  $\mathbf{b}_f, \mathbf{b}_i, \mathbf{b}_o, \mathbf{b}_c$
- initialization:  $\mathbf{c}_0, \mathbf{h}_0$



## Illustration<sup>2</sup>

Input from below, output above, memory (yellow) in the middle.



<sup>2</sup>Illustration by Lobacheva Julia.

# Comments

- Architecture excluded repetitive multiplication of state by the same matrix  $W_s$  (which cause vanishing and exploding gradients)
- Gating mechanisms allow for gradients related to  $c_t$  to stay high across very long time ranges.
- It's recommended to initialize  $\mathbf{b}_f = \mathbf{1}$ 
  - so initially neural net tries to remember everything